

Negotiated Local Plan Repair for Multi-Robot Warehouse Pathfinding

Sreenitya Thatikunta Arpita Deshmukh Prajna Agrawal
B23CS1072 B23CM1007 B23CS1054

1 Problem and environment

n robots move on a 4-connected grid. In each step a robot moves to a neighbouring free cell or waits. Vertex conflicts (two robots in one cell) and edge conflicts (swaps) are forbidden. Each robot has 3 pickup→delivery tasks and then returns to its own dock. The map is an 18×33 Kiva-style warehouse: 5×5 shelf blocks of 2×5 cells separated by one-cell aisles. It has 344 free cells, 150 pickup cells, 14 stations and 58 dead-end docks.

There are three kinds of disruption: **breakdowns** (5–20 steps, or permanent; a permanently broken robot’s tasks are auctioned to robots in range), **blockages** (10–40 steps, or permanent) and **emergency tasks** (a new pickup–delivery task that goes first, with a deadline of the ideal delivery time + 3 steps). Events are generated from the seed, so all strategies face the same schedule. The *dynamic obstacle density* ρ is the number of blockages divided by the number of free cells. By default a filter keeps every instance solvable (Section 4 turns it off). Robots within Manhattan distance $R = 6$ exchange messages (ALERT, REQUEST, PROPOSE, COMMIT, ...), and every message is counted.

2 Method

Initial plan. Prioritized planning: robots are planned one at a time (longest task chain first) with ST-A* over states $(x, y, \text{goal index}, t)$. Each path is written into a space-time reservation table (vertex, edge and parking reservations), and later robots plan around it.

Local repair. When a disruption hits, the robots whose plans it invalidates are repaired one at a time, trying the options that change the fewest plans first (Figure 1). For robot a , let ℓ^* be its robot-free path and $d(p) = |p| - |\ell^*|$ the delay of a path p . A coalition S of neighbours that replan jointly with a (with a planned first) is scored by

$$\text{score}(S) = d(p_a) + \sum_{m \in S} \Delta_m + \lambda|S| + \text{DL}(p_a), \quad (1)$$

where Δ_m is member m ’s added completion time, $\lambda = 3$ is the price of changing one more robot’s plan, and DL adds a penalty of 10 if an emergency deadline is missed. A coalition is accepted only if it scores better than going solo.

- **Tier 1, solo:** replan a alone with all other plans fixed. Accept if the delay is at most $\delta = 4$ (0 for emergency robots).
- **Tiers 2–3, negotiation** within R , then $2R$. Candidates are the movable robots that block ℓ^* . Coalitions of size 1, 2, 3 are tried in order, and the search stops at the first size that beats going solo. A member that cannot dodge may recruit its own blockers (a *chain*, at most 2 hops and 6 robots), scored as one tree.
- **Tier 3b, local group replan:** for urgent robots (emergency, or no solo path), prioritized ST-A* for the at most 8 robots within $2R$ that cross a ’s path.
- **Tier 4, hold:** otherwise a takes its solo detour or, if none exists, waits in place and retries every step. After 10 steps on hold its search radius doubles (expanding ring).




Figure 1: Repair tiers. A robot moves to the next tier only if the current one cannot beat going solo, so the number of robots touched grows only when needed.

We compare **local** (the method above), **local-flat** (our first version, without chains, Tier 3b or deadlines), **solo** (Tier 1 only, an ablation) and **full**, which re-runs prioritized ST-A* for every robot after every disruption. **full** is exactly what the assignment forbids, and serves as a reference.

3 Experiments and results

Each setting was run for 10 seeds with every strategy on identical instances. Defaults are 3% density (20% of blockages permanent), 2 breakdowns and 2 emergencies per run. We varied the number of robots (5–50), the density (0–15% with 20 robots), the alter penalty ($\lambda = 0$ –12), and ran single-disruption runs of each type. **Altered** counts the robots whose future plan changed because of a disruption (including cascades). Every executed step is checked for collisions. All 1,700 runs had zero collisions, and all completed every task except one **solo** run (Section 4).

Table 1: Number of robots (means over 10 seeds, 12 disruptions per run). SoC = sum of task completion times.

n	strategy	SoC	altered /dis.	emerg. on time	ms /dis.	n	strategy	SoC	altered /dis.	emerg. on time	ms /dis.
20	local	2612	2.06	65%	126	50	local	6958	5.88	35%	2270
	local-flat	2608	2.11	45%	81		local-flat	6974	5.75	15%	1102
	solo	2622	1.93	30%	12		solo	7089	4.62	15%	67
	full	2611	9.22	100%	41		full	7201	39.33	100%	590




Figure 2: Plans changed per disruption vs. number of robots (left) and dynamic obstacle density (right, 20 robots). Bands are 95% confidence intervals.

Plans changed. Local repair changes about one robot beyond those a disruption hits directly. Full replanning changes 9–39 robots per disruption, nearly the whole fleet, and is not faster: at 50 robots its SoC is *worse* (7201 vs 6958), because replanning everyone reshuffles priorities and delays robots that were not affected. In single-disruption runs, a temporary breakdown changes 2.2 plans (exactly the robots hit), a temporary blockage 3.1, and a permanent breakdown 13.5 (mostly the robots hit directly plus the winners of its tasks), against 22–29 for full replanning.

Total time and scaling. **local** has the lowest or near-lowest SoC at every size. Negotiation matters more as the fleet grows: **solo** needs 131 more steps at 50 robots. Compared with **local-flat**, the final version lowers SoC by 7–19 steps at 30–50 robots and makes more emergencies on time (35%

vs 15% at 50 robots), changing about as many plans. The cost is about twice the repair time, which reaches 2.3s per disruption at 50 robots. From 0% to 15% density, the extra time grows from 143 to 503 steps while plans changed per disruption stay between 2.1 and 2.7. Raising λ from 0 to 12 trades a few extra steps for fewer altered robots; at 40 robots our default $\lambda = 3$ beats the first version at any λ on both measures.




Figure 3: Left: plans changed to handle a single disruption, by type (10–50 robots pooled). Right: extra time steps caused by disruptions vs. number of robots.




Figure 4: Plans changed vs. extra time as the alter penalty λ varies (ours vs. our first version).

4 Settings where agents fail

In the main sweeps robots rarely fail, partly because the generator filters out disruptions that make a task impossible. A **stress study** (760 runs, 10 seeds, all strategies) switches the filter off and pushes the environment further. A run *fails* if a task is undone when it ends (every robot done, or nothing moved for 150 steps). Each stuck robot is labelled *goal cut off* (its goal is unreachable, so the instance is unsolvable) or *deadlocked* (every goal reachable, but stuck).





Figure 5: Share of tasks completed in the stress settings (10 seeds per point).

Table 2: Where agents fail: runs failed out of 10 and tasks completed; stuck robots per run under `local` (cut off / deadlocked).

setting	value	local (ours)			local-flat	full
		failed	completed	stuck: cut / dead	completed	completed
filter off, half of events permanent	$\rho = 0\%$	3/10	96.9%	0.5 / 0.2	96.9%	96.9%
	$\rho = 5\%$	7/10	91.0%	1.5 / 0.3	91.0%	90.8%
	$\rho = 10\%$	10/10	80.6%	2.3 / 3.1	79.2%	71.3%
	$\rho = 20\%$	10/10	26.8%	9.9 / 8.8	27.4%	24.8%
one perm. breakdown, filter off	20 robots	2/10	97.5%	0.4 / 0.2	97.5%	97.5%
	40 robots	4/10	96.7%	0.4 / 1.0	96.7%	96.7%
very dense, filter on	$\rho = 25\%$	0/10	100%	—	96.3%	100%
	$\rho = 30\%$	1/10	92.9%	0.0 / 2.0	92.7%	86.1%
crowded 12×27 map	30 robots	1/10	92.1%	0.0 / 2.9	92.1%	100%
	40 robots	3/10	73.9%	0.0 / 11.9	90.1%	94.9%
	44 robots	0/10	100%	—	93.4%	100%
comm. radius, 40 robots	$R = 1, 2, 3, 6$	0/10	100%	—	100%	100%



"1 blockage / breakdown": a single permanent event with the solvability filters off. unsafe x%: filters off at density x.

Figure 6: Why robots fail to finish under our method: goal cut off (red) or deadlocked although every goal is reachable (orange).

Goals cut off. With the filter off, a permanent blockage or a robot that breaks down for good can seal off a pickup cell, a station or an aisle segment (Figure 7). No repair can finish those tasks, so all strategies fail alike: completion drops from 97% to 27% as density rises to 20%. A single permanent breakdown already makes 2–4 of 10 runs fail, because a dead robot in a one-cell aisle is a wall.

Knock-on deadlock. A robot whose goal is cut off never gives up its task and holds its cell forever, which can block other robots' only route. At 10% density, 2.3 robots per run are cut off and another

Failure: gridlock without negotiation (solo ablation), 50 robots, 1 permanent breakdown, seed 5
t = 189: 26/150 tasks done; 49 robots stuck (0 goal cut off, 49 deadlocked)




Figure 8: Gridlock without negotiation (**solo**): after one permanent breakdown all 49 remaining robots (red) wait on each other. Our method completes all 150 tasks on the same instance.

5 Graphical demonstration

Simulations are rendered with matplotlib: an animated GIF of a 20-robot run with mixed disruptions (repaired robots and paths highlighted), an interactive replay in the browser, before/after snapshots of each disruption type (old plan dashed, repaired plan solid), GIFs of the failure cases above, and a live replay window.

6 Limitations

Prioritized planning is incomplete, so neither the initial plan nor joint replans are guaranteed to succeed. Robots never release a task whose goal is cut off. Hold cascades on crowded maps need an escalation to a larger regional replan. Emergency deadlines are soft. Repair takes seconds at 50 robots in Python. Communication is assumed instant and lossless.

Code, results and demos: <https://github.com/SreenityaThatikunta/warehouse-plan-repair>